

```

import pandas as pd

# Create the dataset
data = {
    'Pregnancies': [6, 1, 8, 1, 0, 5, 3, 10, 2, 8],
    'Glucose': [148, 85, 183, 89, 137, 116, 78, 115, 197, 125],
    'BloodPressure': [72, 66, 64, 66, 40, 74, 50, 0, 70, 96],
    'SkinThickness': [35, 29, 0, 23, 35, 0, 32, 0, 45, 0],
    'Insulin': [0, 0, 0, 94, 168, 0, 88, 0, 543, 0],
    'BMI': [33.6, 26.6, 23.3, 28.1, 43.1, 25.6, 31.6, 0, 30.5, 0],
    'DiabetesPedigreeFunction': [0.627, 0.351, 0.672, 0.167, 2.288, 0.201, 0.248, 0.134, 0.158, 0.235],
    'Age': [50, 31, 32, 21, 33, 30, 26, 29, 53, 54],
    'Outcome': [1, 0, 1, 0, 1, 0, 1, 0, 1, 1]
}

# Create a DataFrame
df = pd.DataFrame(data)

# Save it to CSV
df.to_csv("pima_indian_diabetes.csv", index=False)

print("Dataset saved as pima_indian_diabetes.csv")

```

➡ Dataset saved as pima_indian_diabetes.csv

```

# Load the dataset
df = pd.read_csv('pima_indian_diabetes.csv')

# Print the first few rows of the dataset to check it
print(df.head())

```

➡

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```

# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset (from the previously saved CSV file)
df = pd.read_csv('pima_indian_diabetes.csv')

```

```

df = pd.read_csv( pima_indian_diabetes.csv )

# Display basic info about the dataset (optional)
print(df.info())

# Split the data into features (X) and target (y)
X = df.drop('Outcome', axis=1) # Features (all columns except 'Outcome')
y = df['Outcome'] # Target variable (1 = Diabetic, 0 = Non-diabetic)

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the features (important for models like Random Forest)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize the Random Forest Classifier model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the model
rf_model.fit(X_train_scaled, y_train)

# Predict on the test set
y_pred = rf_model.predict(X_test_scaled)

# Evaluate the model's accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')

# Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print(f'Confusion Matrix:\n{conf_matrix}')

# Classification Report (precision, recall, F1-score)
print(f'Classification Report:\n{classification_report(y_test, y_pred)}')

# Plot the confusion matrix
sns.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues", xticklabels=["No Diabetes", "Diabetes"],
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# Feature Importances (to understand which features are important for prediction)
feature_importances = rf_model.feature_importances_
features = X.columns
feature_importance_df = pd.DataFrame({'Feature': features, 'Importance': feature_importances})
feature_importance_df = feature_importance_df.sort_values(by='Importance', ascending=False)

# Display feature importance
print(f'Feature Importance:\n{feature_importance_df}')

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   Pregnancies            10 non-null    int64
1   Glucose                10 non-null    int64
2   BloodPressure          10 non-null    int64
3   SkinThickness          10 non-null    int64
4   Insulin                10 non-null    int64
5   BMI                   10 non-null    float64
6   DiabetesPedigreeFunction 10 non-null    float64
7   Age                   10 non-null    int64
8   Outcome                10 non-null    int64

```

dtypes: float64(2), int64(7)

memory usage: 852.0 bytes

None

Accuracy: 50.00%

Confusion Matrix:

```
[[0 1]
```

```
[0 1]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	1
1	0.50	1.00	0.67	1
accuracy			0.50	2
macro avg	0.25	0.50	0.33	2
weighted avg	0.25	0.50	0.33	2

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is zero. The F1 score and weighted average of


_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is zero. The F1 score and weighted average of

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is zero. The F1 score and weighted average of

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))



Start coding or [generate](#) with AI.

Feature Importance:

Feature Importance

1	Glucose	0.229967
6	DiabetesPedigreeFunction	0.186719
5	BMI	0.155099
7	Age	0.147550
2	BloodPressure	0.128113
0	Pregnancies	0.081667
3	SkinThickness	0.036204
4	Insulin	0.034683